

# Representations and Use of Sparse/Structured Matrices

## 1. Introduction to Sparse Matrices

A **sparse matrix** is a matrix that contains a high proportion of zero elements. Storing and computing with a full dense representation is inefficient in terms of memory and computation. Instead, we use **sparse matrix formats** that store only nonzero elements.

### Example: A Sparse Matrix

```
using SparseArrays  
  
A = sparse([1, 2, 3, 4], [4, 3, 2, 1], [5, 3, 2, 1], 4, 4)  
println(A)
```

```
sparse([4, 3, 2, 1], [1, 2, 3, 4], [1, 2, 3, 5], 4, 4)
```

`sparse(row_indices, col_indices, values, num_rows, num_cols)` where:

- `row_indices = [1, 2, 3, 4]` → The row positions of nonzero values.
- `col_indices = [4, 3, 2, 1]` → The column positions of nonzero values.
- `values = [5, 3, 2, 1]` → The actual values at those positions.
- `num_rows = 4` → The total number of rows in the matrix.

Row / Col	1	2	3	4
-----------	---	---	---	---

- `num_cols = 4` → The total number of columns in the matrix.

Row / Col	1	2	3	4
<b>1</b>	0	0	0	5
<b>2</b>	0	0	3	0
<b>3</b>	0	2	0	0
<b>4</b>	1	0	0	0

The output: `sparse([4, 3, 2, 1], [1, 2, 3, 4], [1, 2, 3, 5], 4, 4)` where: - Rows and columns are reordered for storage efficiency. - The format `sparse(row_indices, col_ptrs, values, m, n)` means: - `row_indices` → Positions of nonzero elements (sorted by column). - `col_ptrs` → The starting index of each column in `row_indices`. - `values` → The corresponding nonzero values.

## 2. Sparse Matrix Representations

Sparse matrices are stored in **compressed formats** to optimize memory usage.

### (a) Compressed Sparse Row (CSR)

- Stores nonzero values row-wise.
- Efficient for row-based operations.

### (b) Compressed Sparse Column (CSC)

- Stores nonzero values column-wise.
- Efficient for column-based operations.

### (c) Coordinate List (COO)

- Stores (row, col, value) triplets.
- Simple to construct but inefficient for arithmetic operations.

```
rows = [1, 2, 3, 4]
cols = [4, 3, 2, 1]
values = [5, 3, 2, 1]
A = sparse(rows, cols, values)
println(A)
```

```
sparse([4, 3, 2, 1], [1, 2, 3, 4], [1, 2, 3, 5], 4, 4)
```

### 3. Special Types of Structured Matrices

In addition to sparsity, some matrices have **mathematical structures** that allow for optimized storage and computation.

#### (a) Diagonal Matrices

- Only the main diagonal has nonzero values.
- Storage complexity:  $O(n)$ .

#### (b) Triangular Matrices

- Upper or lower triangular matrices contain nonzero elements only above or below the diagonal.

#### (c) Symmetric Matrices

- Only one half of the matrix is stored.

#### (d) Toeplitz Matrices

- Constant diagonals, common in signal processing.

#### (e) Band Matrices

- Nonzero elements are confined to a diagonal band.

```
using SparseArrays, LinearAlgebra # Import LinearAlgebra for tril and triu

# Lower triangular sparse matrix
L = tril(sprand(5, 5, 0.3))

# Upper triangular sparse matrix
U = triu(sprand(5, 5, 0.3))

println("Lower Triangular Sparse Matrix:")
println(L)
```

```
println("Upper Triangular Sparse Matrix:")
println(U)

# Symmetric Sparse Matrix
A = sprand(5, 5, 0.3)
A_sym = A + A' # Ensure symmetry

# Toeplitz Matrix
using ToeplitzMatrices
T = Toeplitz([1, 2, 3], [1, 4, 5])
println(T)

# Band Matrix
B = spdiags(-1 => [1, 1, 1], 0 => [4, 4, 4], 1 => [2, 2, 2])
```

Lower Triangular Sparse Matrix:

```
sparse([1, 2, 4, 2, 4, 4, 5], [1, 1, 1, 2, 2, 3, 4], [0.796955628263455, 0.27090577143668015
```

Upper Triangular Sparse Matrix:

```
sparse([1, 2, 1, 1], [2, 2, 3, 4], [0.9157743284168461, 0.002547265286128919, 0.744946491635
[1 4 5; 2 1 4; 3 2 1]
```

4×4 SparseMatrixCSC{Int64, Int64} with 9 stored entries:

```
4  2
1  4  2
   1  4  2
     1
```

## 4. Representations and Use of Sparse/Structured Matrices with Performance Analysis

### 1. Performance Metrics in Julia

Julia provides built-in tools for benchmarking: - `@time`: Basic timing and memory usage. - `@btime` (from `BenchmarkTools`): More accurate benchmarking. - `@allocated`: Measures memory allocation.

```
using SparseArrays, LinearAlgebra, BenchmarkTools
```

## 2. Dense vs. Sparse Matrix Operations

### (a) Creating a Large Dense vs. Sparse Matrix

We compare memory usage when creating a large random matrix:

```
n = 5000 # Matrix size
```

```
@time dense_matrix = rand(n, n)
```

```
@time sparse_matrix = sprand(n, n, 0.01) # 1% nonzero entries
```

```
0.029740 seconds (16.50 k allocations: 191.530 MiB, 14.10% gc time, 14.26% compilation time)
```

```
0.001264 seconds (10 allocations: 3.868 MiB)
```

5000×5000 SparseMatrixCSC{Float64, Int64} with 249471 stored entries:

### Expected Outcome:

- Dense matrices allocate significantly more memory ( $O(n^2)$ ).
- Sparse matrices store only nonzero values, saving space ( $O(n)$ ).

## (b) Matrix-Vector Multiplication

We compare the efficiency of multiplying a dense vs. sparse matrix with a vector:

```
using BenchmarkTools # Import the package

v = rand(n)

@btime dense_matrix * v
@btime sparse_matrix * v
```

```
1.708 ms (3 allocations: 39.19 KiB)
89.916 s (3 allocations: 39.19 KiB)
```

5000-element Vector{Float64}:

```
17.640553482070384
11.254848100389935
14.138314864666349
12.083067684779547
13.664857446125604
 6.762099284339201
11.742508778120495
 8.804010538292463
 9.08203108745566
 9.820964828825309

10.165766223424583
11.0202507112116
16.24321306922888
12.156443706282394
15.840870316074714
 9.832429791126831
13.62619264883956
 7.886635622599546
16.50638263201445
```

### Expected Outcome:

Sparse multiplication should be much faster when nonzero elements are a small fraction.

### (c) Solving a Linear System $Ax = b$

Solving  $Ax = b$  for a dense and sparse matrix:

```
b = rand(n)

@btime x_dense = dense_matrix \ b
@btime x_sparse = sparse_matrix \ b
```

```
213.598 ms (9 allocations: 190.81 MiB)
465.263 ms (95 allocations: 760.61 MiB)
```

5000-element Vector{Float64}:

```
-13.510436188911536
-18.20763312012673
 9.683269136278183
 4.129184612275613
 9.876873483985023
-1.9960062191354098
12.92452684419674
-4.0402341845972005
 6.692382409266598
-5.736136183502744

-0.3610529681919809
-30.481223042489866
 9.497345513785309
-8.63326630618668
-9.644787347756553
-6.544122862199334
-18.70316341133351
 0.39778294976645867
-3.5750162352512187
```

### Why is the Sparse Solve Slower?

1. Direct Solvers ( $\backslash$ ) are Optimized for Dense Matrices
  - The backslash operator `\` in Julia uses LU decomposition for dense matrices, which is highly optimized with BLAS and LAPACK.
  - However, for sparse matrices, `\` defaults to a direct solver (SuperLU), which is not always the best choice.

## 2. Sparse Matrices Have Fill-in During Factorization

- When performing LU factorization on a sparse matrix, many zero elements become nonzero (fill-in), increasing memory and computational costs.

## 3. Sparse Matrices Work Best with Iterative Solvers

- Instead of using \, iterative solvers like Conjugate Gradient (CG) or GMRES are often much faster and more memory efficient.

### Understanding Preconditioning in Iterative Solvers

Preconditioning is a technique to **improve the efficiency** of iterative solvers for solving sparse linear systems of the form:

$$Ax = b$$

where: - A is a sparse (often large and ill-conditioned) matrix. - x is the unknown vector. - b is the right-hand-side vector.

### Why Use Preconditioning?

Without preconditioning, iterative solvers (like GMRES, CG, BiCGSTAB) may: - Converge very slowly (or not at all) for poorly conditioned matrices. - Require many iterations, making them computationally expensive.

Preconditioning transforms the original system into an equivalent but easier system, reducing the condition number and accelerating convergence.

### Using Preconditioning for the previous example:

```
using IterativeSolvers
@btime x_sparse_cg = cg(sparse_matrix, b)
@btime x_sparse_gmres = gmres(sparse_matrix, b)

using IncompleteLU

# Compute ILU preconditioner
M = ilu(sparse_matrix)

# Solve Ax = b using GMRES with preconditioner
@btime x_sparse_preconditioned = gmres(sparse_matrix, b, Pl=M)
```



```
536.487 ms (20 allocations: 157.39 KiB)
674.862 ms (526 allocations: 958.03 KiB)
9.171 s (528 allocations: 958.30 KiB)
```

5000-element Vector{Float64}:

NaN  
NaN  
NaN  
NaN  
NaN  
NaN  
NaN  
NaN  
NaN  
NaN  
NaN

NaN  
NaN  
NaN  
NaN  
NaN  
NaN  
NaN  
NaN  
NaN  
NaN

## Why is GMRES + ILU Slower Despite Preconditioning?

### 1. ILU Factorization Overhead

- `ILU(ilu(sparse_matrix))` must first compute an approximate factorization before GMRES can use it.
- This setup takes time and allocates memory.

### 2. Sparse Matrix Fill-in

- ILU may increase the number of nonzero elements (“fill-in”), making GMRES less efficient.

### 3. Preconditioner Effectiveness

- If ILU doesn’t sufficiently reduce condition number, GMRES still needs many iterations, slowing it down.

### 3. Optimized Sparse Matrix Operations

#### Dense Matrix Storage

- A dense matrix of size  $n \times n$  requires storing every element.
- If  $n = 5000$ , a dense matrix (Float64) requires:  $5000 \times 5000 \times 8 = 200MB$  (since Float64 takes 8 bytes per entry)

#### Sparse Matrix Storage

- A sparse matrix only stores:
  - Nonzero values
  - Row indices
  - Column pointers
- If the matrix has only 1% nonzero elements ( $\sim 250,000$  values in a  $5000 \times 5000$  matrix), then:
  - Values:  $250,000 \times 8 \text{ bytes} = \sim 2MB$
  - Row indices:  $250,000 \times 4 \text{ bytes} = \sim 1MB$
  - Column pointers:  $5000 \times 4 \text{ bytes} = \sim 20KB$
  - Total storage 3MB instead of 200MB for a dense matrix.

#### (a) Using Iterative Solvers

Instead of `\`, we can use Conjugate Gradient (CG), which is optimized for sparse systems.

```
using IterativeSolvers

@btime x_sparse_cg = cg(sparse_matrix, b)
```

```
528.806 ms (20 allocations: 157.39 KiB)
```

```
5000-element Vector{Float64}:
 1.5808999111691338e19
 2.2374004830958744e19
 6.033497027581915e18
 1.814394428119684e18
-1.1381000939018973e19
 4.697671071592917e18
 1.5032525774910353e19
 1.3547694672325046e19
```

```

1.4148754230082796e19
-1.374882334402571e19

-2.773095803071125e19
1.612985179462213e19
2.0470425661963543e19
-2.2219028592460313e19
5.372747133585884e18
-8.598024368037902e18
1.4816899735230458e19
1.6105223695506008e19
7.839996407776828e18

```

### Expected Outcome:

- CG converges much faster for large sparse matrices.

### (b) Memory Comparison: Allocations in Dense vs. Sparse Matrices

We measure memory usage during operations:

```

using SparseArrays, BenchmarkTools

n = 5000 # Matrix size
dense_matrix = rand(n, n) # Dense matrix
sparse_matrix = sprand(n, n, 0.01) # 1% nonzero elements

v = rand(n) # Correctly sized vector

# Measure memory allocations
mem_dense = @allocated dense_matrix * v
mem_sparse = @allocated sparse_matrix * v

println("Memory allocated for Dense Matrix-Vector Multiplication: ", mem_dense, " bytes")
println("Memory allocated for Sparse Matrix-Vector Multiplication: ", mem_sparse, " bytes")
println("Memory used by Dense Matrix: ", Base.summarysize(dense_matrix), " bytes")
println("Memory used by Sparse Matrix: ", Base.summarysize(sparse_matrix), " bytes")

```

```

Memory allocated for Dense Matrix-Vector Multiplication: 40128 bytes
Memory allocated for Sparse Matrix-Vector Multiplication: 40128 bytes
Memory used by Dense Matrix: 200000048 bytes
Memory used by Sparse Matrix: 4065280 bytes

```

### Expected Outcome:

- Both operations take up the same amount memory, but sparse matrix takes up less memory usage.

## 4. Special Matrix Efficiency Comparisons

We evaluate computation speed using structured matrices.

### (a) Diagonal Matrices

Diagonal matrices can be stored efficiently using `Diagonal()` instead of a full matrix.

```
D = Diagonal(rand(n))

@btime D * v # Fast multiplication
@btime Matrix(D) * v # Slow multiplication
```

```
1.742 s (3 allocations: 39.19 KiB)
13.073 ms (6 allocations: 190.77 MiB)
```

5000-element Vector{Float64}:

```
0.041345892433626415
0.026443473726081204
0.005863695455556071
0.1746351513784327
0.12492310495959255
0.19738377211492739
0.13558622424036654
0.45036121374812565
0.015273953694571185
0.32248909747405863

0.1072372049499046
0.0067917837949993485
0.02358368309050058
0.41893366880262856
0.0019218217571370941
0.10030100025098668
0.2166356284377104
0.15988748250034288
0.07264189059738102
```

### Expected Outcome:

- Direct `Diagonal()` multiplication is much faster than using a full matrix.

### (b) Toeplitz Matrices

Toeplitz matrices can be stored using a single row and column.

```
using ToeplitzMatrices, BenchmarkTools

n = 1000
first_element = rand() # Generate a shared first element

T = Toeplitz([first_element; rand(n-1)], [first_element; rand(n-1)]) # Ensure first element

v = rand(n)

@btime T * v
@btime Matrix{T} * v # Dense conversion is inefficient
```

```
173.833 s (21 allocations: 102.34 KiB)
879.917 s (6 allocations: 7.64 MiB)
```

1000-element Vector{Float64}:

```
260.65794294077415
258.19301016896424
257.90858166021195
259.21902908196955
260.0424786168712
259.62516053784674
261.56730352439234
261.0060182332099
260.0930734925788
262.1612864721572
```

```
255.68209254890408
259.35091370634115
259.4088405155768
255.7584476415949
255.2138648796962
255.61483039287876
```

258.7918474443171  
258.55222645950255  
257.9787794818542

### Expected Outcome:

Using the specialized Toeplitz representation avoids unnecessary memory usage.

## 5. Using GPUs for Sparse and Structured Matrices in Julia

### 1. Why Use GPUs for Sparse Matrices?

Advantages of GPUs: - Parallelism: GPUs have thousands of cores, making them ideal for matrix-vector and matrix-matrix multiplications. - High Memory Bandwidth: GPUs access memory faster than CPUs for structured operations. - Efficient for Iterative Solvers: CG, GMRES, BiCGSTAB are well-suited for GPUs.

Common Applications: - Computational Physics & Engineering: Solving large sparse linear systems from Finite Element Methods (FEM). - Machine Learning: Speeding up graph-based models (e.g., PageRank, Graph Neural Networks). - Fluid Dynamics & PDEs: Solving Poisson equations with sparse linear solvers.

### 2. GPU-Based Special Matrices

Julia's CuSparse library provides GPU-accelerated sparse matrix operations.

#### (a) Sparse Matrix in GPU

Convert a CPU sparse matrix to GPU using `CuSparseMatrixCSC`, then compare time usage between CPU and GPU:

```
using BenchmarkTools, CUDA

function benchmark_spmv(n=10^4, density=0.01)
    # CPU setup
    A_cpu = sprand(n, n, density)
    x_cpu = rand(n)

    # GPU setup
    A_gpu = CuSparseMatrixCSC(A_cpu)
    x_gpu = CuArray(x_cpu)
```

```

# Warm-up
A_cpu * x_cpu
A_gpu * x_gpu

# Benchmark
cpu_time = @belapsed $A_cpu * $x_cpu
gpu_time = @belapsed CUDA.@sync $A_gpu * $x_gpu

println("CPU: $(cpu_time*1e3) ms | GPU: $(gpu_time*1e3) ms")
end

benchmark_spmv(10^4, 0.01)

```

```

julia> benchmark_spmv(10^4, 0.01)
CPU: 2.213315 ms | GPU: 0.096481 ms

```

## (b) Structured Matrices Example (Block Diagonal)

Use GPU-accelerated solvers from `CUDA.jl`'s `CuSPARSE` library:

```

using SparseArrays, CUDA, BenchmarkTools

function block_diagonal_spmv(n_blocks=100, block_size=10)
    nnz_per_block = block_size^2
    total_nnz = nnz_per_block * n_blocks
    I = Int[]
    J = Int[]
    V = Float64[]
    sizehint!(I, total_nnz)
    sizehint!(J, total_nnz)
    sizehint!(V, total_nnz)

    for b in 0:(n_blocks-1)
        offset = b * block_size
        block = rand(block_size, block_size) # Dense block
        for i in 1:block_size, j in 1:block_size
            push!(I, offset + i)
            push!(J, offset + j)
            push!(V, block[i, j])
        end
    end
end

```

```

end

N = n_blocks * block_size
A_cpu = sparse(I, J, V, N, N)
x_cpu = rand(N)

A_gpu = CuSparseMatrixCSC(A_cpu) # Convert to GPU sparse format
x_gpu = CuArray(x_cpu)

println("CPU (Sparse Block Diagonal):")
cpu_time = @belapsed $A_cpu * $x_cpu
println("Time: $(cpu_time * 1e3) ms")

println("GPU (Sparse Block Diagonal):")
CUDA.@sync A_gpu * x_gpu # Warm-up to remove initialization overhead
gpu_time = @belapsed CUDA.@sync $A_gpu * $x_gpu
println("Time: $(gpu_time * 1e3) ms")

return cpu_time, gpu_time
end

block_diagonal_spmv(100, 10)
block_diagonal_spmv(1000, 100)

```

```

julia> block_diagonal_spmv(1000, 100) # Increase problem size
CPU (Sparse Block Diagonal):
Time: 14.399187000000001 ms
GPU (Sparse Block Diagonal):
Time: 0.352362 ms
(0.014399187, 0.000352362)

julia> block_diagonal_spmv(100, 10)
CPU (Sparse Block Diagonal):
Time: 0.013184 ms
GPU (Sparse Block Diagonal):
Time: 0.056926 ms
(1.3184e-5, 5.6926e-5)

```

Figure 1: image.png

Troubleshooting with the smaller matrix sizes:



- CPU: 0.014 ms
  - CPU is surprisingly faster than GPU, which is unusual for large matrices but expected for small problems due to kernel launch overhead.
- GPU: 0.056 ms. 4x slower than CPU. This is likely because:
  - The matrix is too small, and GPU has a high overhead for launching kernels.
  - Kernel execution time is overshadowed by memory transfer between CPU & GPU.

### 3. Performance Considerations for GPU Sparse Matrices in Julia

When working with sparse matrices on GPUs, performance depends on matrix format and kernel optimization techniques. Below are example codes for each key concept:

#### (a) Performance Considerations for GPU Sparse Matrices in Julia

When working with sparse matrices on GPUs, performance depends on matrix format and kernel optimization techniques. Below are example codes for each key concept.

#### 1. Compressed Sparse Column (CSC) & Compressed Sparse Row (CSR)

- CSC (Column-wise storage): Efficient for column slicing, but not ideal for SpMV on GPUs.
- CSR (Row-wise storage): Better for GPU performance, especially with CUDA CUSPARSE.

```
# Load Packages
using CUDA, SparseArrays, LinearAlgebra, BenchmarkTools

# Generate a Large Sparse Matrix
# Problem size
n = 10^5 # Matrix size
density = 0.001 # Fraction of nonzero elements
# Generate a random sparse matrix on CPU (ensure Float32 for GPU compatibility)
A = sprand(Float32, n, n, density)
A = A + A' + n * I # Make it symmetric positive definite

# Move Data to GPU
# Move sparse matrix to GPU in CSR format
dA_CSR = CuSparseMatrixCSR(A) # Ensure Float32 type
# Move vector to GPU with matching type (Float32)
```

```

x = CUDA.rand(Float32, n) # Input vector on GPU
y = CUDA.zeros(Float32, n) # Output vector on GPU (preallocated)

# Perform GPU Sparse Matrix-Vector Multiplication
CUDA.CUSPARSE.mv!('N', 1.0f0, dA_CSR, x, 0.0f0, y, 'O')

```

**Fix Explanation:** - 'N': No transpose ('T' for transpose).

- 1.0f0: Float32 scaling factor for Ax (use 1.0 for Float64).
- dA\_CSR: The sparse GPU matrix (must be CuSparseMatrixCSR{Float32}).
- x: The input vector on GPU (must be CuArray{Float32}).
- 0.0f0: Float32 scaling factor for y.
- y: The preallocated output vector.
- 'O': Optimized CUDA algorithm for SpMV.

```

#Benchmark CPU vs. GPU Performance
#CPU
x_cpu = rand(Float32, n) # CPU input vector
y_cpu = zeros(Float32, n) # CPU output vector
println("CPU Sparse Matrix-Vector Multiplication Benchmark:")
@btime mul!($y_cpu, $A, $x_cpu) # CPU SpMV

#GPU
println("GPU Sparse Matrix-Vector Multiplication Benchmark:")
@btime CUDA.CUSPARSE.mv!('N', 1.0f0, $dA_CSR, $x, 0.0f0, $y, 'O')

```

```

julia> println("CPU Sparse Matrix-Vector Multiplication Benchmark:")
CPU Sparse Matrix-Vector Multiplication Benchmark:

julia> @btime mul!($y_cpu, $A, $x_cpu) # CPU SpMV
63.979 ms (0 allocations: 0 bytes)
100000-element Vector{Float32}:

julia> println("GPU Sparse Matrix-Vector Multiplication Benchmark:")
GPU Sparse Matrix-Vector Multiplication Benchmark:

julia> @btime CUDA.CUSPARSE.mv!('N', 1.0f0, $dA_CSR, $x, 0.0f0, $y, 'O')
42.871 μs (86 allocations: 1.78 KiB)
100000-element CuArray{Float32, 1, CUDA.DeviceMemory}:

```

## 2. Block Compressed Sparse Row (BSR) Format

1. What is BSR Format? Block Compressed Sparse Row (BSR) is an extension of CSR, where nonzero elements are stored in small dense blocks instead of individual elements. Best for structured sparsity, where nonzeros naturally cluster into blocks. More efficient than CSR when matrix structure fits block patterns.

2. Correct BSR Implementation in CUDA.jl CUDA.jl does support BSR format, but we need to: - Start with a CPU sparse matrix (`SparseMatrixCSC`).

- Convert it to a GPU sparse format (`CuSparseMatrixCSC`).
- Convert it to `CuSparseMatrixBSR` (GPU BSR format).
- Use `CUDA.CUSPARSE.mv!` for efficient SpMV.

**Move Sparse Matrix to GPU and Convert to BSR Format:**

```
# Convert CPU sparse matrix to GPU sparse matrix (CSC first)
dA_CSC = CuSparseMatrixCSC(A)

# Convert CSC to BSR format (use block size = 4)
block_size = 4
dA_BSR = CuSparseMatrixBSR(dA_CSC, block_size)
```

**Fix Explanation:**

- Convert *A* to `CuSparseMatrixCSC` first, because `CuSparseMatrixBSR` only accepts GPU sparse matrices.
- Then, convert `CuSparseMatrixCSC` to `CuSparseMatrixBSR`.

**Move Dense Vector to GPU:**

```
x = CUDA.rand{Float32, n} # Input vector on GPU
y_BSR = CUDA.zeros{Float32, n} # Output vector on GPU (preallocated)
```

**Perform GPU Sparse Matrix-Vector Multiplication:**

```
CUDA.CUSPARSE.mv!('N', 1.0f0, dA_BSR, x, 0.0f0, y_BSR, 'O')
```

**Fix Explanations:**

- `mv!` performs efficient SpMV with BSR format.

Arguments: - 'N': No transpose ('T' for transpose).

- 1.0f0: Float32 scaling factor for Ax (use 1.0 for Float64).
- dA\_BSR: BSR format matrix on GPU.
- x: Input vector on GPU.
- 0.0f0: Float32 scaling factor for y\_BSR.
- y\_BSR: Preallocated output vector.
- '0': Optimized CUDA algorithm for SpMV.

Benchmark CPU vs. GPU Performance:

```
#CPU
x_cpu = rand(Float32, n) # CPU input vector
y_cpu = zeros(Float32, n) # CPU output vector
println("CPU Sparse Matrix-Vector Multiplication Benchmark:")
@btime mul!($y_cpu, $A, $x_cpu) # CPU SpMV

#GPU
println("GPU Sparse Matrix-Vector Multiplication (BSR) Benchmark:")
@btime CUDA.CUSPARSE.mv!('N', 1.0f0, $dA_BSR, $x, 0.0f0, $y_BSR, '0')
```

```
julia> println("CPU Sparse Matrix-Vector Multiplication Benchmark:")
CPU Sparse Matrix-Vector Multiplication Benchmark:

julia> @btime mul!($y_cpu, $A, $x_cpu) # CPU SpMV
63.855 ms (0 allocations: 0 bytes)

julia> println("GPU Sparse Matrix-Vector Multiplication (BSR) Benchmark:")
GPU Sparse Matrix-Vector Multiplication (BSR) Benchmark:

julia> @btime CUDA.CUSPARSE.mv!('N', 1.0f0, $dA_BSR, $x, 0.0f0, $y_BSR, '0')
11.312 μs (35 allocations: 656 bytes)
100000-element CuArray{Float32, 1, CUDA.DeviceMemory}:
```

### 3. Hybrid Formats (ELL, HYB)

- ELL (Ellpack-Itpack): Stores fixed-length rows (good for structured sparsity).
- HYB (Hybrid): Combines ELL + COO for improved performance.
- Hybrid formats (HYB) optimize SpMV when sparsity patterns vary.

## (b) Kernel Optimizations for GPU Sparse Matrices

Beyond choosing the right format, optimizing GPU kernels is crucial.

### 1. Thread Blocking: Grouping Nonzeros for Better Memory Access

- SpMV on GPU can suffer from non-coalesced memory access.
- Thread blocking groups nonzero elements to optimize access patterns.

### 2. Shared Memory: Caching Frequently Accessed Data

- Using shared memory reduces redundant global memory access.
- This is advanced but highly effective for improving SpMV.

#### Summary:

**Sparse Matrix Formats** - CSC/CSR: General-purpose but may not optimize for GPU.

- Block Compressed Sparse Row (BSR): Better for structured sparsity.
- Hybrid Formats: Custom formats for specific problems.

**Kernel Optimization** - Thread Blocking: Group nonzeros for coalesced memory access.

- Shared Memory: Cache frequently accessed data (advanced).

## 6. Summary of Calculation Choices

Method	Time Complexity	Memory Efficiency
<b>Dense Matrix Multiplication (with Vector)</b>	$O(n^2)$	High memory usage
<b>Sparse Matrix Multiplication (with Vector)</b>	$O(n)$ (depends on sparsity)	Low memory usage
<b>Direct Solve (\)</b>	$O(n^3)$ for dense	Fast for small matrices
<b>Iterative Solve (CG)</b>	$O(n)$ for sparse	Best for large matrices
<b>Special Matrices (Diagonal, Toeplitz, Band)</b>	Optimized operations	Highly efficient

## 7. Conclusion

- **Sparse matrices** reduce memory and speed up computations.
- **Iterative solvers** like **Conjugate Gradient** outperform direct solvers for large sparse systems.
- **Structured matrices** (diagonal, banded, Toeplitz) allow for optimized operations.

## 8. Possible Real-World Applications

### (a) Finite Element Methods

- Problem: Solve  $Ku = f$ , where  $K$  is a sparse stiffness matrix.
- GPU Advantage: Assembly of  $K$  and iterative solvers (e.g., conjugate gradient).

### (b) Graph Neural Networks (GNNs)

- Sparse Adjacency Matrices: Accelerate message-passing steps on GPUs.

### (c) Quantum Lattice Models

- Hamiltonians: Sparse matrix diagonalization with ARPACK (GPU port).